

IMPORTING PACKAGES

```
In [73]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import train_test_split

from sklearn import preprocessing

from sklearn.ensemble import RandomForestRegressor
```

LOADING DATA

```
In [74]: raw_data = pd.read_csv('C:/Users/ronin/Google Drive/IAS ML MODEL/BIG DATA/big_dat.csv')
# print(raw_data.info())
```

```
In [75]: raw_data.columns = ['REF'] + [''] * (len(raw_data.columns)-1) #making REF column of raw_data
raw_data.head()
```

```
Out[75]:
```

	REF									
0	1.3008	0.005880	0.005925	0.005982	0.006050	0.006133	0.006233	0.006352	0.006493	0.006653
1	1.3012	0.005875	0.005920	0.005975	0.006042	0.006124	0.006222	0.006340	0.006480	0.006640
2	1.3016	0.005870	0.005914	0.005968	0.006035	0.006115	0.006212	0.006328	0.006466	0.006627
3	1.3020	0.005865	0.005908	0.005962	0.006027	0.006106	0.006201	0.006316	0.006453	0.006613
4	1.3024	0.005860	0.005903	0.005955	0.006020	0.006097	0.006191	0.006304	0.006439	0.006598

5 rows × 100 columns



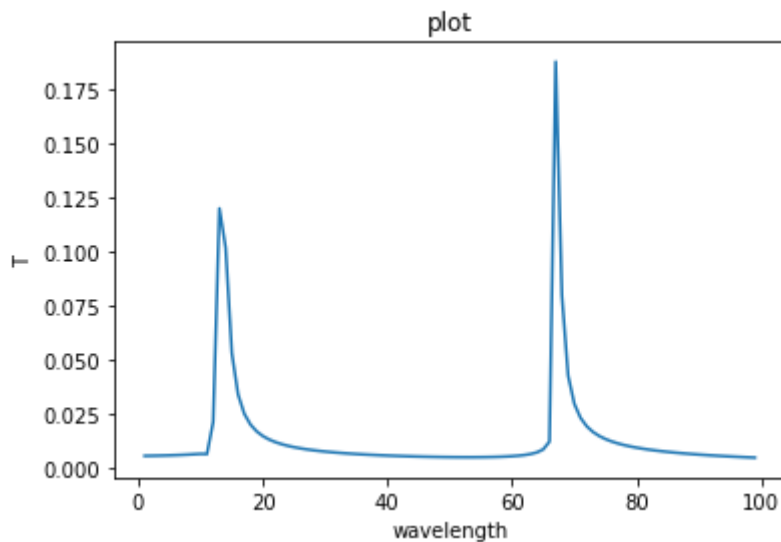
PLOTTING A ROW

```
In [76]: # r2 = raw_data.iloc[:,1:]
x = list(range(1, 100))
y = raw_data.iloc[1,1:]
plt.plot(x, y)

# naming the x axis
plt.xlabel('wavelength')
# naming the y axis
plt.ylabel('T')

# giving a title to my graph
plt.title('plot')

# function to show the plot
plt.show()
```



```
In [77]: y = raw_data['REF'] #Y is our label vector
raw_data.drop(['REF'], axis=1, inplace=True)
print(y)
```

```
0      1.3008
1      1.3012
2      1.3016
3      1.3020
4      1.3024
...
494    1.4984
495    1.4988
496    1.4992
497    1.4996
498    1.5000
Name: REF, Length: 499, dtype: float64
```

```
In [78]: raw_data.columns =list(range(1,100)) #adding column names as features
raw_data.columns
raw_data.head()
```

```
Out[78]:
```

	1	2	3	4	5	6	7	8	9	10
0	0.005880	0.005925	0.005982	0.006050	0.006133	0.006233	0.006352	0.006493	0.006653	0.006793
1	0.005875	0.005920	0.005975	0.006042	0.006124	0.006222	0.006340	0.006480	0.006640	0.006794
2	0.005870	0.005914	0.005968	0.006035	0.006115	0.006212	0.006328	0.006466	0.006627	0.006791
3	0.005865	0.005908	0.005962	0.006027	0.006106	0.006201	0.006316	0.006453	0.006613	0.006783
4	0.005860	0.005903	0.005955	0.006020	0.006097	0.006191	0.006304	0.006439	0.006598	0.006773

5 rows × 99 columns



SPLITTING TO TRAIN AND TEST DATA

```
In [173... X_train, X_test, y_train, y_test = train_test_split(raw_data, y, test_size=0.2, rand
raw_data.shape
y_train.shape
```

```
Out[173... (399,)
```

APPLYING RANDOM FOREST REGRESSION MODEL

```
In [174... regr = RandomForestRegressor(n_estimators=3, max_depth=100, max_leaf_nodes=500, min_
m1 = regr.fit(X_train,y_train)
m1.score(X_test,y_test)
```

Out[174... 0.999855925686978

CROSS VALIDATED

```
In [193... from sklearn import model_selection
from sklearn import metrics
from sklearn.model_selection import KFold
model = RandomForestRegressor(n_estimators=3, max_depth=3, max_leaf_nodes=500, min_s
#cv = model_selection.KFold(n_splits=3)

kf = KFold(n_splits=10, random_state=4200, shuffle=True) # shuffling returns good
# kf = KFold(n_splits=2,shuffle=False) #not shuffling returns worse results hence m
kf.get_n_splits(raw_data)

for train_index, test_index in kf.split(raw_data):
#     print("TRAIN:", train_index, "TEST:", test_index)
    X_train, X_test = raw_data.iloc[train_index], raw_data.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    # For training, fit() is used
    m1 = model.fit(X_train, y_train)

    # Default metric is R2 for regression, which can be accessed by score()

    print(m1.score(X_test, y_test))
```

```
0.9868795740176696
0.988476777398903
0.9897716049132753
0.9909167860482109
0.9880294490403972
0.9892151955067838
0.9850462643639619
0.9882777814134562
0.9918181075865021
0.9879701223180762
```

```
In [157... regr.decision_path(X_test)
```

Out[157... (<49x1525 sparse matrix of type '<class 'numpy.int64'>'
with 1309 stored elements in Compressed Sparse Row format>,
array([0, 511, 1032, 1525], dtype=int32))

VISUALIZING TREE

```
In [196... from sklearn.tree import plot_tree

fig = plt.figure(figsize=(15, 10))
plot_tree(regr.estimators_[1],
          feature_names=X_test.columns,
```

```
        filled=True, impurity=True,  
        rounded=True)  
# it takes average of datapoints at leaf and returns as predicted value
```

MAKE PREDICTIONS FROM TEST SET

```
In [189... y_pred = (m1.predict(X_test))
```

CALCULATE RMS ERROR

```
In [183... from sklearn.metrics import mean_squared_error  
  
rms = mean_squared_error(y_test, y_pred, squared=True)  
print(rms)
```

8.001334668498421e-07

PREDICTATED VALUES

```
In [184... y_pred
```

```
Out[184... array([1.30775259, 1.31904848, 1.31904848, 1.33292905, 1.33593571,  
        1.33698333, 1.33997619, 1.3433404 , 1.3449      , 1.35554      ,  
        1.35650667, 1.35780667, 1.35935      , 1.3648309 , 1.36592296,  
        1.36592296, 1.36973667, 1.37268571, 1.38671      , 1.39540889,  
        1.39854444, 1.40069286, 1.40231952, 1.40668571, 1.40885333,  
        1.41520238, 1.42107333, 1.42350963, 1.43501429, 1.43701984,  
        1.44334444, 1.4456      , 1.44693333, 1.44776667, 1.44895714,  
        1.45101524, 1.45802222, 1.45802222, 1.46078889, 1.47305905,  
        1.47435667, 1.47956429, 1.48095476, 1.48498535, 1.48624672,  
        1.49468202, 1.49749841, 1.49749841, 1.49749841])
```

```
In [185... y_test
```

```
Out[185... 17      1.3076  
44      1.3184  
47      1.3196  
78      1.3320  
87      1.3356  
91      1.3372  
94      1.3384  
109     1.3444  
111     1.3452  
138     1.3560  
139     1.3564  
142     1.3576  
147     1.3596  
160     1.3648  
163     1.3660  
167     1.3676  
168     1.3680  
180     1.3728  
215     1.3868  
236     1.3952  
244     1.3984  
249     1.4004  
259     1.4044  
260     1.4048  
269     1.4084  
285     1.4148  
302     1.4216  
305     1.4228  
339     1.4364
```

```

340    1.4368
357    1.4436
360    1.4448
366    1.4472
369    1.4484
372    1.4496
373    1.4500
392    1.4576
396    1.4592
397    1.4596
431    1.4732
434    1.4744
449    1.4804
450    1.4808
458    1.4840
462    1.4856
483    1.4940
493    1.4980
495    1.4988
498    1.5000
Name: REF, dtype: float64


```

In [190]...

```

plt.scatter(y_test,y_pred, c='coral')
plt.xlabel('True Values ')
plt.ylabel('Predictions ')
# plt.axis('equal')
# plt.axis('square')
plt.show()

```




In [191]...

```

x = np.linspace(0,90,y_test.shape[0])
plt.scatter(x, y_pred, color = 'red') # red shows predicted
plt.scatter(x, y_test, color = 'blue') #blue is actual
plt.show()
# hence perfect overlap shows good prediction

```



In []: